



Homework 03

NAME: Rosetta Wang

STUDENT ID: 3034171690

Numpy Introduction

1a) Create two numpy arrays (a and b). a should be all integers between 25-34 (inclusive), and b should be ten evenly spaced numbers between 1-6. Print all the results below:

- i) Cube (i.e. raise to the power of 3) all the elements in both arrays (element-wise)
- ii) Add both the cubed arrays (e.g., $[1,2] + [3,4] = [4,6]$)
- iii) Sum the elements with even indices of the added array.
- iv) Take the square root of the added array (element-wise square root)___

```
In [1]: import numpy as np
a = np.array(range(25, 35))
b = np.linspace(1.0, 6.0, num=10)
print(a)
print(b)
a **= 3
b **= 3
print(a)
print(b)

[25 26 27 28 29 30 31 32 33 34]
[1.          1.55555556 2.11111111 2.66666667 3.22222222 3.77777778
 4.33333333 4.88888889 5.44444444 6.          ]
[15625 17576 19683 21952 24389 27000 29791 32768 35937 39304]
[ 1.          3.76406036  9.40877915 18.96296296 33.45541838
 53.91495199 81.37037037 116.85048011 161.38408779 216.          ]
```

1b) Append b to a, reshape the appended array so that it is a 4x5, 2d array and store the results in a variable called m. Print m.

```
In [17]: m = np.array([a, b], np.int32).reshape(4, 5)
print(m)
```

```
[[15625 17576 19683 21952 24389]
 [27000 29791 32768 35937 39304]
 [    1     3     9    18    33]
 [   53    81   116   161   216]]
```

1c) Extract the third and the fourth column of the m matrix. Store the resulting 4x2 matrix in a new variable called m2. Print m2.

```
In [13]: # I assume you mean to store a 2x4 matrix because you can't take the dot product
m2 = np.column_stack([m[:,2], m[:,3]])
print(m2)
m2 = np.array([m[:,2], m[:,3]])
print(m2)
```

```
[[19683 21952]
 [32768 35937]
 [    9    18]
 [   116   161]]
[[19683 32768     9    116]
 [21952 35937    18    161]]
```

1d) Take the dot product of m2 and m store the results in a matrix called m3. Print m3. Note that Dot product of two matrices $A.B = A^T B$

```
In [14]: m3 = np.dot(m2, m, out = None)
print(m3)
```

```
[[1192289032 1322149319 1461175850 1609683670 1767987512]
 [1313307551 1456440614 1609683670 1773384518 1947890546]]
```

1e) Round the m3 matrix to three decimal points. Store the result in place and print the new m3.

```
In [16]: m3 = np.around(m3, decimals=3)
print(m3)
```

```
[[1192289032 1322149319 1461175850 1609683670 1767987512]
 [1313307551 1456440614 1609683670 1773384518 1947890546]]
```

1f) Sort the m3 array so that the highest value is at the bottom right and the lowest value is at the top left. Print the sorted m3 array.

```
In [18]: print(np.sort(m3))
```

```
[[1192289032 1322149319 1461175850 1609683670 1767987512]
 [1313307551 1456440614 1609683670 1773384518 1947890546]]
```

NumPy and Masks

2a) create an array called 'f' where the values are cosine(x) for x from 0 to pi with 50 equally spaced values in f

- print f
- use a 'mask' and print an array that is True when $f \geq 1/2$ and False when $f < 1/2$
- create and print an array sequence that has only those values where $f \geq 1/2$

```
In [21]: f = np.cos(np.linspace(0, np.pi, num=50))
print(f)
mask = f >= 1/2
print(f[mask])
```

```
[ 1.          0.99794539  0.99179001  0.98155916  0.96729486  0.94905575
 0.92691676  0.90096887  0.8713187   0.8380881   0.80141362  0.76144596
 0.71834935  0.67230089  0.6234898   0.57211666  0.51839257  0.46253829
 0.40478334  0.34536505  0.28452759  0.22252093  0.1595999   0.09602303
 0.03205158 -0.03205158 -0.09602303 -0.1595999   -0.22252093 -0.28452759
-0.34536505 -0.40478334 -0.46253829 -0.51839257 -0.57211666 -0.6234898
-0.67230089 -0.71834935 -0.76144596 -0.80141362 -0.8380881   -0.8713187
-0.90096887 -0.92691676 -0.94905575 -0.96729486 -0.98155916 -0.99179001
-0.99794539 -1.          ]
[1.          0.99794539  0.99179001  0.98155916  0.96729486  0.94905575
 0.92691676  0.90096887  0.8713187   0.8380881   0.80141362  0.76144596
 0.71834935  0.67230089  0.6234898   0.57211666  0.51839257]
```

NumPy and 2 Variable Prediction

Let 'x' be the number of miles a person drives per day and 'y' be the dollars spent on buying car fuel (per day).

We have created 2 numpy arrays each of size 100 that represent x and y.

x (number of miles) ranges from 1 to 10 with a uniform noise of (0,1/2)

y (money spent in dollars) will be from 1 to 20 with a uniform noise (0,1)

```
In [29]: # seed the random number generator with a fixed value
import numpy as np
np.random.seed(500)

x=np.linspace(1,10,100)+ np.random.uniform(low=0,high=.5,size=100)
y=np.linspace(1,20,100)+ np.random.uniform(low=0,high=1,size=100)
print ('x = ',x)
print ('y= ',y)
```

```
x = [ 1.34683976  1.12176759  1.51512398  1.55233174  1.40619168  1.6507
5498
 1.79399331  1.80243817  1.89844195  2.00100023  2.3344038  2.22424872
 2.24914511  2.36268477  2.49808849  2.8212704  2.68452475  2.68229427
 3.09511169  2.95703884  3.09047742  3.2544361  3.41541904  3.40886375
 3.50672677  3.74960644  3.64861355  3.7721462  3.56368566  4.01092701
 4.15630694  4.06088549  4.02517179  4.25169402  4.15897504  4.26835333
 4.32520644  4.48563164  4.78490721  4.84614839  4.96698768  5.18754259
 5.29582013  5.32097781  5.0674106  5.47601124  5.46852704  5.64537452
 5.49642807  5.89755027  5.68548923  5.76276141  5.94613234  6.18135713
 5.96522091  6.0275473  6.54290191  6.4991329  6.74003765  6.81809807
 6.50611821  6.91538752  7.01250925  6.89905417  7.31314433  7.20472297
 7.1043621  7.48199528  7.58957227  7.61744354  7.6991707  7.85436822
 8.03510784  7.80787781  8.22410224  7.99366248  8.40581097  8.28913792
 8.45971515  8.54227144  8.6906456  8.61856507  8.83489887  8.66309658
 8.94837987  9.20890222  8.9614749  8.92608294  9.13231416  9.55889896
 9.61488451  9.54252979  9.42015491  9.90952569  10.00659591  10.02504265
10.07330937  9.93489915  10.0892334  10.36509991]
y= [ 1.6635012  2.0214592  2.10816052  2.26016496  1.96287558  2.95546
35
 3.02881887  3.33565296  2.75465779  3.4250107  3.39670148  3.39377767
 3.78503343  4.38293049  4.32963586  4.03925039  4.73691868  4.30098399
 4.8416329  4.78175957  4.99765787  5.31746817  5.76844671  5.93723749
 5.72811642  6.70973615  6.68143367  6.57482731  7.17737603  7.54863252
 7.30221419  7.3202573  7.78023884  7.91133365  8.2765417  8.69203281
 8.78219865  8.45897546  8.89094715  8.81719921  8.87106971  9.66192562
 9.4020625  9.85990783  9.60359778  10.07386266  10.6957995  10.66721916
11.18256285  10.57431836  11.46744716  10.94398916  11.26445259  12.09754828
12.11988037  12.121557  12.17613693  12.43750193  13.00912372  12.86407194
13.24640866  12.76120085  13.11723062  14.07841099  14.19821707  14.27289001
14.30624942  14.63060835  14.2770918  15.0744923  14.45261619  15.11897313
15.2378667  15.27203124  15.32491892  16.01095271  15.71250558  16.29488506
16.70618934  16.56555394  16.42379457  17.18144744  17.13813976  17.69613625
17.37763019  17.90942839  17.90343733  18.01951169  18.35727914  18.16841269
18.61813748  18.66062754  18.81217983  19.44995194  19.7213867  19.71966726
19.78961904  19.64385088  20.69719809  20.07974319]
```

3a) Find Expected value of x and the expected value of y

```
In [31]: print(np.mean(x), np.mean(y))

5.782532541587923 11.012981683344968
```

3b) Find variance of distributions of x and y

```
In [33]: np.var(x)
```

```
Out[33]: 7.03332752947585
```

```
In [34]: np.var(y)
```

```
Out[34]: 30.113903575509635
```

3c) Find co-variance of x and y.

```
In [36]: np.cov(x,y)[0][1]
```

```
Out[36]: 14.657743832803439
```

3d) Assuming that number of dollars spent in car fuel is only dependant on the miles driven, by a linear relationship.

Write code that uses a linear predictor to calculate a predicted value of y for each x ie $y_{\text{predicted}} = f(x) = y_0 + mx$.

```
In [38]: X = np.column_stack([np.ones(len(x)),x])
W = np.linalg.inv(X.T.dot(X)).dot(X.T).dot(y)
```

```
In [39]: y_predicted = lambda x: W[0]+W[1]*x
```

3e) Predict y for each value in x, put the error into an array called y_error

```
In [44]: y_error = y_predicted(x) - x
print(y_error)
```

```
[ 0.5144174  0.27512051  0.6933373  0.73289662  0.5775204  0.83754028
 0.98983137  0.99880996  1.10088125  1.20992128  1.5643962  1.44727924
 1.47374909  1.59446454  1.73842587  2.08203311  1.93664503  1.9342736
 2.37318137  2.22638221  2.36825421  2.54257519  2.71373237  2.70676278
 2.81081081  3.06904065  2.96166494  3.09300495  2.87136955  3.34687688
 3.50144492  3.39999276  3.36202194  3.60286053  3.50428164  3.62057272
 3.68101899  3.85158317  4.16977317  4.23488484  4.36336127  4.5978554
 4.71297616  4.73972383  4.47013098  4.90455547  4.89659827  5.08462264
 4.92626266  5.35273607  5.12727262  5.20942846  5.40438857  5.65447973
 5.42468354  5.49094901  6.0388744  5.99233916  6.24846926  6.33146315
 5.99976594  6.43490137  6.53816125  6.41753574  6.85779669  6.74252303
 6.63581927  7.03731915  7.15169508  7.18132783  7.2682202  7.43322632
 7.62538881  7.38379768  7.8263278  7.58132407  8.01952064  7.89547377
 8.0768316  8.16460551  8.32235703  8.24572096  8.4757272  8.29306689
 8.59638029  8.87336784  8.61030294  8.57267417  8.79193935  9.24548463
 9.3050085  9.2280809  9.09797185  9.61827121  9.72147633  9.74108893
 9.79240614  9.64524829  9.80933658 10.10263805]
```

```
In [ ]: # I assume you mean "put the error"
```

3f) Write code that calculates the root mean square error(RMSE), that is root of average of y-error squared

```
In [47]: np.sqrt(np.mean(y_error ** 2))
```

```
Out[47]: 5.94205703941763
```

```
In [ ]:
```